

CITS2006:
Defensive Cybersecurity
Portfolio
Part 2

Submitted by:

Kelly Valencia

24540356

<https://github.com/Kelly-VI/Defensive-Cybersecurity-Portfolio>

Table of Self-Assessment

Activity	Description	Marks
B1	Discover 5 unique weak/vulnerable security implementations	2
B2	Discover 5 unique strong security implementations	2
B3	Discover 3 proactive security implementations in practice	2
B4	Participate in 3 in-class activities in labs (facilitators will administer such activities)	2
B5	Attend 2 cybersecurity related talks/seminars	2
B6	Attend 2 cyber ethics/law related talks/seminars	1
B7	Participate in CS/DS/cybersecurity clubs' activity	2
B8	Participate in a hackathon	0
B9	Participate in an industry-related cybersecurity event	2
B10	Participate in a cybersecurity activity as part of job/internship/volunteering etc	0
B11	Talk to 2 cybersecurity experts from the industry and find out their latest projects	2
B12	Discover two bias cases when using a generative AI system	2
B13	Perform a jailbreak attack on a generative AI assistant (controlled test only)	2
B14	Teach your friends about cybersecurity topic of your choice	2
B15	Teach an elderly person about cybersecurity topic of your choice	2
B16	Survey the current state-of-the-art solutions in cybersecurity	2
B17	Implement one of the current state-of-the-art solutions and evaluate it	0
B18	Contribute to an open-source project related to cybersecurity	2
B19	Find and fix a vulnerability from a GitHub project	2
B20	Enhance the security of a GitHub project	2
B21	Design and implement a cybersecurity learning activity	2
B22	Enhance the cybersecurity of a website from your community	2
B23	Test an intrusion detection system and discuss its effectiveness	2
B24	Design and implement access control of your choice	2
B25	Design and implement a threat intelligence module of your choice	0
B26	Help another student in this unit struggling to understand/learn a cybersecurity concept	2
B27	Apply a learned concept in this unit to a real-world application/problem/environment	2
B28	Produce a cyber safety flyer for (choose 1): elders, high school students, CEOs, Uni students	2
B29	Find a CVE in this year and fix it using three different generative AI systems (e.g. ChatGPT, Gemini), comparing the consistency	2
B30	Generate an AI-created image, applying an imperceptible watermark on it and then perform an image-to-image regeneration or editing process to make sure the watermark is detectable---the watermark survives	0

Done: 25/25

Part 2 Activities

(Due end of Week 12, Friday 11.59pm AWST).

Choose any 25 activities from the following list to complete and submit.

B1. Discover 5 unique weak/vulnerable security implementations

For this activity, I will be basing of my research from [CWE \(Common Weakness Enumeration\)](#), which is a category system for hardware and software weaknesses and vulnerabilities. When searching through a specific CWE ID, they will sometimes give Selected Observed Examples of the CWE through a [CVE \(Common Vulnerabilities and Exposures\)](#), which is a system that provides a reference method for publicly known information-security vulnerabilities and exposures.

1. [CWE-327: Use of a Broken or Risky Cryptographic Algorithm](#)

This CWE means that a system or product uses a broken or risky cryptographic algorithm or protocol. If the cryptographic algorithm is not strong enough, the confidentiality, integrity, and repudiation of sensitive data may be compromised. An (old) example of this is [CVE-2008-3775](#), where Folder Lock 5.9.5 and earlier uses ROT-25 for the password, which is a weak encryption as it allows local administrators to obtain sensitive information by reading and decrypting the QualityControl_pack registry value. ROT-25 is a Caesar cipher, which is one of the simplest encryption techniques as it is simply a substitution cipher and easily broken.

2. [CWE-330: Use of Insufficiently Random Values](#)

This CWE means that a system uses insufficiently random numbers or values in a security context, such as a protection, authorization, or authentication mechanism, that depends on unpredictable numbers. If the mechanism relies on random values to restrict access to a sensitive resource, then the resource being protected could be accessed by guessing the “random number” since it was taken from a small enough sample size, violating confidentiality and access control since now attackers can gain access to resources from other user’s accounts. An example of this is [CVE-2009-2158](#), where account-recover.php in TorrentTrader Classic 1.09 chooses random passwords from an insufficiently large set, which makes it easier for remote attackers to obtain a password via a brute-force attack.

3. [CWE-1188: Initialization of a Resource with an Insecure Default](#)

This CWE uses a system that initializes or sets a resource with a default that is intended to be changed by the product's installer, administrator, or maintainer, but the default is

not secure. Default values/configuration are often used since developers assume that administrators can and will change the default values to something more secure, so they leave the product open and make it easy to use as possible out-of-the-box for users to adjust. However, the problem comes when the default is insecure and the administrator does not change the settings. An example of this is [CVE-2022-36349](#), where there was an insecure default variable initialization in BIOS firmware for some Intel(R) NUC Boards and Intel(R) NUC Kits before version MYi30060 that may allow an authenticated user to potentially enable denial of service via local access.

4. [CWE-1104: Use of Unmaintained Third Party Components](#)

This CWE uses a system that relies on third-party components that are not actively supported or maintained by the original developer or a trusted proxy for the original developer. Relying on unmaintained components makes it difficult or impossible to fix significant bugs and vulnerabilities, can render code obsolete, and undermine security by complicating maintenance and increasing the risk of new vulnerabilities. An example of this is [CVE-2025-40906](#), where Perl module for BSON serialization includes a component, a bundled libbson 1.1.7, that reached end-of-life as of August 13, 2020 (approximately 5 years previously as of the discovery of this CVE) and is no longer supported. This distribution has several vulnerabilities, such as CVE-2017-14227, CVE-2018-16790, CVE-2023-0437, CVE-2024-6381, CVE-2024-6383, and CVE-2025-0755.

5. [CWE-284: Improper Access Control](#)

This CWE uses a system that does not restrict or incorrectly restricts access to a resource from an unauthorized actor. Access control involves authentication, authorization, and accountability. If any of the access control mechanisms fail or is not applied, attackers can compromise the security of the product by gaining privileges, reading sensitive information, executing commands, evading detection, etc. Weaknesses in access control is mostly introduced due to incorrect specification (e.g. privileges, permissions, ownership) specified for either the user or the resource, and errors in the enforcement of access control. An example of this is [CVE-2021-37415](#), where Zoho ManageEngine ServiceDesk Plus before 11302 is vulnerable to authentication bypass that allows a few REST-API URLs without authentication.

B2. Discover 5 unique strong security implementations.

1. **Zero Trust Architecture (ZTA)**

- https://en.wikipedia.org/wiki/Zero_trust_architecture

ZTA is a design and implementation strategy of IT systems that holds the principle that users and devices should not be trusted by default, even if they are connected to a

privileged network or were previously verified. This is implemented by establishing identity verification, validating device compliance prior to granting access, and ensuring least privilege access to only explicitly-authorized resources. The approach advocates mutual authentication, including checking the identity and integrity of users and devices without respect to location, and providing access to applications and services based on the confidence of user and device identity and device status in combination with user authentication. ZTA has been adopted into the Government and Healthcare sector. It is a strong security implementation because even if attackers breach one system, lateral movement becomes extremely difficult.

2. Hardware Security Modules (HSM)

- https://en.wikipedia.org/wiki/Hardware_security_module
- <https://www.cyber.gov.au/glossary/hardware-security-module>

A HSM is a physical computing device that safeguards and manages cryptographic keys and provides cryptographic processing (encryption and decryption functions) for digital signatures, strong authentication, etc. They are commonly deployed in Public Key Infrastructure, digital identity solutions and payment systems, often coming in the form of a plug-in card or an external device that attaches directly to a computer or network server as they contain one or more secure cryptoprocessor chips that prevent tampering and bus probing. They have features that provide tamper evidence such as visible signs of tampering or logging and alerting, or tamper resistance which makes tampering difficult without making the HSM inoperable, or tamper responsiveness such as deleting keys upon tamper detection. It is a strong security implementation because attackers cannot easily extract encryption keys even after compromising servers.

3. WebAuthn/FIDO2 Passwordless Authentication

- <https://frontegg.com/guides/passwordless-authentication-with-fido2-and-webauthn>
- <https://fidoalliance.org/specifications/>

FIDO2 is comprised of the W3C Web Authentication specification and corresponding Client-to-Authenticator Protocols (CTAP) from the FIDO Alliance. It supports passwordless, second-factor and multi-factor user experiences with embedded (or bound) authenticators (such as biometrics or PINs) or external (or roaming) authenticators (such as FIDO Security Keys, mobile devices, wearables, etc.) and uses public-key cryptography.

WebAuthn (Web Authentication) is a W3C standard that defines a browser-based API for creating and using public-key credentials and enables support for FIDO Authentication. It allows users to authenticate using devices like security keys, smartphones, or built-in biometric systems (e.g., fingerprint readers), without relying on passwords.

It is a strong security implementation since it eliminates most password theft and phishing attacks. Using public-key cryptography eliminates the storage of passwords on servers, which reduces risk of breaches and credential theft. They also use site-specific key pairs, ensuring that credentials are not reusable across different domains, which makes it impossible for attackers to use phishing techniques to access sensitive information.

4. Runtime Application Self-Protection (RASP)

- <https://www.crowdstrike.com/en-us/cybersecurity-101/cloud-security/runtime-application-self-protection-rasp/>
- https://en.wikipedia.org/wiki/Runtime_application_self-protection

RASP is a security technology that incorporates security functionality within software applications to prevent malicious attacks while the application is running. RASP does this by installing sensors embedded within the software to monitor and control application execution during runtime. They combine the sensor data with contextual information (e.g. code, application logic, settings and configuration, runtime data, control flow, inputs) to address specific vulnerabilities that exist within each piece of software as it identifies high-risk activity or active breaches, and make accurate, timely decisions about how to mitigate such events and stop threats automatically and in real time. This is a strong security implementation because it offers automated, real-time protection at the application level, which can protect application attacks such as Zero Day attacks, Cross Site Scripting (XSS), SQL injections, and DoS/DDoS.

5. Confidential Computing

- <https://www.ibm.com/think/topics/confidential-computing>
- https://en.wikipedia.org/wiki/Confidential_computing

Confidential Computing is a cloud computing technology that protects data during processing. Before data can be processed by an application, it must be unencrypted in memory, leaving it vulnerable. So Cloud Computing isolated the data by using a hardware-based trusted execution environment (TEE) (a secure enclave within a CPU) which is secured using embedded encryption keys, where embedded attestation mechanisms ensure that the keys are accessible to authorized application code only as TEE will deny access to the keys and cancel computation to any malware or unauthorized code attempting to access the keys. This means sensitive data can remain protected in memory until the application tells the TEE to decrypt it for processing. While the data is decrypted throughout the entire computation process, it is invisible to the operating system, the hypervisor in a virtual machine (VM), to other compute stack resources and to the cloud service provider and its employees. This is a strong security implementation because it protects sensitive data even while in use and protects it from all other entities except the authorized programming code itself.

B3. Discover 3 proactive security implementations in practice.

1. Address Space Layout Randomization (ASLR)

- https://en.wikipedia.org/wiki/Address_space_layout_randomization
- <https://web.mit.edu/ha22286/www/papers/Encyclopedia23.pdf>

ASLR is a computer security technique involved in preventing exploitation of memory corruption vulnerabilities by randomly arranging the address space positions of key data areas of a process, including the base of the executable and the positions of the stack, heap, and libraries. This is done in order to prevent an attacker from reliably redirecting code execution to a particularly exploited function in memory. ASLR is part of Moving Target Defenses (MTD), which are techniques that seek to enhance the resiliency of a computer system by making it more dynamic, random, and/or diverse. ASLR has been running on most major desktop and mobile operating systems.

2. Security Information and Event Management (SIEM)

- <https://www.ibm.com/think/topics/siem>
- <https://cybersecuritynews.com/proactive-cybersecurity/>
- <https://www.ibm.com/products/qradar>

SIEM is a security solution that helps organizations recognize and address potential security threats and vulnerabilities before they have a chance to disrupt business operations.

The system collects and aggregates event data from across an organization's IT infrastructure. These event data can include including users, endpoints, applications, networks, cloud workloads, and security tools. The data logs are collected, correlated, and analyzed in real-time. SIEM uses advanced analytics to identify and understand intricate data patterns, event correlation provides insights to quickly locate and mitigate potential threats to business security. Then it consolidates its analysis into a single, central dashboard where security teams monitor activity, triage alerts, identify threats and initiate response or remediation. Due to the automated data collection and analysis that it provides, SIEM can gather and verify compliance data across the entire business infrastructure while also generating real-time compliance reports for PCI-DSS, GDPR, HIPAA, SOX and other compliance standards, reducing the burden of security management and detecting potential violations early so they can be addressed. An example of this is IBM QRadar SIEM.

3. External Attack Surface Management (EASM)

- <https://www.crowdstrike.com/en-us/cybersecurity-101/exposure-management/external-attack-surface-management-easm/>
- <https://www.crowdstrike.com/en-us/platform/exposure-management/easm/>

EASM is the continuous discovery, monitoring, evaluation, prioritization, and remediation of attack vectors of an organization's external attack surface/digital attack surface, which is the sum of an organization's internet-facing assets and the associated attack vectors which can be exploited during an attack. These internet-facing assets include anything from domain names, SSL certificates and protocols to operating

systems, servers, IOT devices and network services, and may be scattered across on-premise, cloud environments and third-party vendors. An example of this is CrowdStrike Falcon Exposure Management.

EASM automatically detects all exposed assets, known or unknown, across all environments and supply chain and view them by category, platform, and service type. They have business context insight (the ability to understand exactly to which business unit, subsidiary or 3rd party vendor each exposed asset belongs and how it is connected to the core network) and uses this to generate security insight and risk prioritization so the most critical tasks can be focuses on.

B4. Participate in 3 in-class activities in labs (facilitators will administer such activities).

In-Class Lab Activity 1 (Thursday 5 March 2026)

Since Week 2 lecture topics was about Cryptography, the class activity in the labs had us form groups and come up with multiple plaintexts, encrypt them with some simple cryptography you can do by hand and get their ciphertexts, and write down the decryption algorithm to decrypt. Then we went around to see what the other groups had written and attempted to decrypt their ciphertexts.

My group had the plain text “Cryptography is hard” and “Lemons are sour”, and our decryption algorithm was to do ROT13 and reverse the text. We solved the other groups’ decryption where one had Morse Code and ROT5, and another had multiple shifts using ROT.

Evidence for this activity will be under directory “B4” with images “In-Class Lab Activity 1.1 - 1.6”.

In-Class Lab Activity 2 (Monday, 9 March 2026)

Since this week’s lecture was about Privacy, the class activity was about Multiple Party Computation (MPC), which is a technique on sharing information to other parties without revealing your own information. In my group, we were told to think of a random number and not tell the others about our number, but somehow get the average of all our numbers combined. We tried following the method shown in the lecture slides (Lecture 03. Privacy Slide 60-63) where we kept a portion of our number to ourselves and split the remaining to the others (some equally shared, some unequal). Then we add up all the numbers we received, including our private portion of the number we kept to get the total, which we then will divide by 5 to get the average. Then we compare our average with the other.

Somehow, we managed to get different results and I think we had some miscalculation or we missed a step along the way since the actual average (once we revealed our numbers to check)


was different from the averages everyone calculated, especially since each person got a different average.

The method the lab facilitator gave was quite simple. The starting person picks a random number and adds it to their number. Then they pass along that number to the others, where each person will then add their number and pass it along, and so on and so forth, until it comes back to that first person, where they would subtract the sum they received with the random number they added in the beginning, which would then give the total needed to find the average.

Evidence for this activity (mostly paper-based calculations and the student names and numbers of our group) will be under directory “B4” with images “In-Class Lab Activity 2.1 – 2.4”.


In-Class Lab Activity 3 (Monday, 16 March 2026)

The topic for this week’s lecture was Access Control, and the lab activity we did was Ruleset-based Access Control, where each of our groups had to come up with 4 rules that had to apply to all the members, then another group has to guess what all our rules were. My evidence for this, as shown below and in the directory “B4” under “In-Class Lab Activity 3.1”, contains the rules my group made, along with the group members’ name and student number. I forgot to take a picture of the rules of the other group we had to guess, but if I remember correctly, one of their rules had to do with a specific amount of laptop Operating Systems (the group was all male, around 5-6 members I think?).



In-Class Lab Activity 4 (Monday, 20 April 2026)

The topic for this week's lecture was Threat Modelling, and our lab activity involved our group going around the CS building to check the security/safety features and give it a score along with justification on why we gave it that score. There were only 5 people in the lab that day, so the groups were small (1-2 people), but we went around the building together and commented on the security. Then we went back to the lab, finished the worksheet, and then discussed with one another why we gave it the score. Each group scored it differently since the scoring was subjective based on what we thought was "secure", so we discussed that having a "criteria" for security would make it more objective and have people score more similarly.




B5. Attend 2 cybersecurity related talks/seminars.

1. Cybersecurity Talk from Vanessa Van Beek

During the CyberQueens x CyberThrive event, there was a cybersecurity talk from Vanessa Van Beek, who shared about her experiences in her career and the ups and downs she experienced to get through where she is now. She is currently the Chief Information Security Officer (CISO) at Fortescue, but before cybersecurity, she revealed that she was studying Law. However, somewhere along the way, she decided to study and go into the field of cybersecurity.

This pivot in her career was definitely not an easy one, and luckily, she had the support from her surroundings to help her out. Some of the analogies she used in this talk were a butterfly/caterpillar cocoon when she was starting out her journey, referencing how the cocoons were like the start of her growth into the field, where she was constantly “consuming” all the information about cybersecurity she could, especially when her interest was still at its peak. Then during her hardships, it was important to have a group that would support you, and this group encompassed 3 roles (using the imagery of dolphins): The attacker, the defender, and the scout.

Throughout that journey, she had trouble with asking for help from others. So, her question at the end of her talk was asking whether everyone liked helping people (many people raised their hands), and whether they liked asking for help (very few people raised their hand). Thus, her call to action was that we should not be afraid to ask for help from others.





2. Week 10 Defensive Cybersecurity Guest Lecture

In Week 10 of our Defensive Cybersecurity lecture, we had a guest lecture hosted by Wellington Muvunduke. He talked about Fraud Detection in the real world, specifically in the banking sector.

He explained how financial fraud and traditional cybersecurity differed in how financial fraud focuses more on human behaviour such as phishing and social engineering in how attackers use them to gain stolen credentials, while traditional cybersecurity focused more on the system and how it is exploited through intrusions or malware.

Then he proceeded to explain the different types of financial crime in the real world: Account Takeover, Authorized Push Payment, Identity Theft, Card Fraud (skimming device copies the magnetic strip of card number onto external machine), Insider Threats, Money Laundering/AML (ask other people to move money in and out). For example, the steps that occur for Account Takeover first includes harvesting credentials, which the attackers will use the access the account and do some reconnaissance, and after that they may do some social engineering to try and bypass MFA to finally transfer funds from the victim's account.

Afterwards, he talked about the different types of scams (Application Fraud, Investment Scams, Romance & Pig Butchering, SIM Swap Attacks, Synthetic Identity Fraud, and Business Email Compromise) and what detection signals analysts look at to detect suspicious behaviour (Technical Signals and Behavioural Signals), while also trying to balance out any False Positives or False Negatives found. To help with fraud detection, he mentioned how AI was being used to do Anomaly Detection, Behavioural Profiling, Real-Time Risk Scoring, and Analyst Decision Support.

All evidence for this activity can be found under the directory /Evidence/B5.

B6. Attend 2 cyber ethics/law related talks/seminars.

I attended the cyber talk/lecture by Larry Huynh on Week 12 (Monday, 18 May 2026). He talked about AI in Defensive Cybersecurity, going from a general view of Artificial Intelligence, then going deeper into the deeper subsets of Machine Learning, Deep Learning, and Large Language Models. During the talk, he showed a demo of Anomaly-based IDS and how different statistics can “draw” the given data differently (t-SNE) and another demo for Malware detection and how they can be converted into 2D grayscale images from their binary representations.

For LLMs, he talked how they were created from billions of datasets/parameters, and they first created a pre-trained model that allowed them to learn how to speak like a human (grammar, syntax, and all that). After that, the model can be fine-tuned and trained into more specific domains. Finally, he talked about using LLMs for CVD and how they were decent for detecting any CWE, but not really for identifying the correct CWE.

In my Evidence/B6 folder I have an image of him doing the presentation as well as my notes I have taken down while Larry was presenting. Unfortunately, I did not manage to get a picture of the signature attendance paper since I had to rush to my group project demo, but I did sign the paper.

B7. Participate in CS/DS/cybersecurity clubs’ activity.

I attended the “Into the Cuckoo’s Egg” cybersecurity tech talk hosted by UCC club and presented by Mark Tearle on 24 April 2026. Mark talked about how the book is written from the perspective of the hero Clifford as he recounts his journey on hunting down the computer hacker Markus who broke into a computer at Lawrence Berkeley National Laboratory (LBNL) as he tracked down the trails left by the hacker and logged his details of his accounts. He told us how some of the actions that Clifford are still relevant to modern cybersecurity, such as reporting the issue and getting the law enforcement early, as well as getting in touch with the right person. Most of the activity that happened here was written under “notes.jpg” in the evidence.

All evidence for this activity will be under the directory /Evidence/B7.

into the cuckoo's egg

a cybersecurity tech talk
presented by mark tearle

week 9 friday
24th apr @ 2pm

at the dawn of the internet era, a 75-cent accounting error at the lawrence berkeley national laboratory draws one systems administrator into a digital spy-hunt.

LIFFS
THE
Cuckoo
E



ucc
university
computer
club

ucc.au/links



B9. Participate in an industry-related cybersecurity event.

The CyberThrive x CyberQueens event hosted by CyberWest and WiTWA had a facilitated networking event, where all the participants were grouped based on their current stage of their cybersecurity career (green for student/currently studying, yellow for early stage in their career, orange for mid stage, and red for executive stage). Me, along with 2 other guys and 3 other

mature aged female students, were grouped with a mentor who asked us questions about what problems have you encountered in your career and what resources have helped you so far in your career. One of the mature aged females in the group mentioned how she had difficulties getting back to her career after a break due to pregnancy, me and the other two guys (both of whom were UWA students) mentioned difficulties in group project and getting internships, while the other 2 mature aged females, who were Murdoch TAFE students, mentioned how they had difficulty in following their lectures because there was no recording where they can rewatch the lecture (which surprised me since UWA had lecture recordings, and we spent some time discussing this topic).

The cybersecurity resources that were mentioned were some websites like W3Schools that could help with the coding aspect of cybersecurity, and the mentor suggested whether we would be interested in a cross-school event where fellow cybersecurity students can connect with and help out one another (or something along those lines).

While talking with the others, I asked about whether they knew some cybersecurity events that were going to take place soon and was recommended to join two events (as shown in evidence B9-rec1 and B9-rec2).

Evidence for this activity can be found under the directory /Evidences/B9.

B11. Talk to 2 cybersecurity experts from the industry and find out their latest projects.

I met these two cybersecurity experts during the CyberThrive x CyberQueens event. All evidence of this (including a picture with them!) can be found under the directory /Evidences/B11.

Caitriona Forde (Cyber Safety & Human Factors Expert):

Forde's latest project has to do with her PhD research into Human Factors in Cyber Safety and how we can bridge the gap between knowing the right thing and actually doing it. For example, we might know that changing passwords regularly, using a password manager, and setting up MFAs are very important things to do in securing our privacy, but people tend to not do it despite knowing this.

Most cybersecurity incidents come from human error, and instead of allocating their budget to fix this human risk, major companies are spending their money on their own technology.

Some tests that were conducted would be sending fake or test emails to employees and see which person will check on them. If they do check on them, then they might need some training in cyber safety awareness.

Biljana R. (Experienced Cyber Governance, Risk & Compliance (GRC) Professional):

Biljana has stated that she works in Governance, Risk & Compliance and critical infrastructure, so most of her projects are confidential. However, what she can share is that when a business

wants to use a new vendor, then said business will come to her company to assess the risk of that vendor. This involves various things, such as how to connect the business to the system, where does the data go, and creating a risk profile of the vendor. Hence, they inform businesses whether or not they can use the vendors after a risk assessment process.

She also mentioned operational technology, which is a physical industrial control system (or according to [IBM](#): the hardware and software that directly monitors, controls, and automates physical processes and infrastructure in industrial settings), and how those systems must be available at all times (one of the CIA triads) as a shutdown of these industrial systems (a hacker getting hold of these systems) can cause safety issues for people. Thus, they have to find ways to ensure that the system keeps running.

B12. Discover two bias cases when using a generative AI system.

The bias cases were taken from this website: <https://www.crescendo.ai/blog/ai-bias-examples-mitigation-guide>.

1. Facebook's Age-Related AI Bias

LinkedIn's AI-driven job recommendation systems faced allegations of perpetuating gender biases as some studies and investigation uncovered that LinkedIn's algorithms favoured male candidates over equally qualified female counterparts via a fairness metric detecting algorithm bias and search algorithm alternative users suggestions, which leads to unequal job recommendations. This made the LinkedIn team build a new AI algorithm to deal with the issue.

2. Apple's Gender-Biased AI for Credit Cards (<https://time.com/5724098/new-york-investigating-goldman-sachs-apple-card/>)

Apple's credit card algorithm went under investigation after complaints circled around the internet which began when someone's husband commented how his card's credit limit was 20 times higher than his wife's, despite her higher credit score and shared financial accounts. The same thing was said by the co-founder of Apple who stated that his credit limit was 10 times higher than his wife's, even though they did not have any separate assets or accounts.


B13. Perform a jailbreak attack on a generative AI assistant (controlled test only).



This controlled test was taken from:

- <https://llm-attacks.org/>
- <https://github.com/llm-attacks/llm-attacks/issues/8> (Recreating the example shown in the github)

The research paper proposed a simple and effective attack method that causes Large Language Models to generate objectionable behaviours. This attack method involves finding a suffix that

maximizes the probability that the model produces an affirmative response to the query to produce objectionable content rather than the model refusing to answer the query. The approach of the paper procudes the adversarial suffixes via a combination of greedy and gradient-based search techniques and is quite transferable to publicly released LLMs. I have recreated the paper's example using the steps given by users in the Issues of their Github, as shown in the images below. The last image shows how if the adversarial siffix is not added at the end of the query, the model, Llama 2 7B Chat, will not answer the query. This is a type of jailbreak prompt called "Prompt Injection", where the attacker embeds malicious directives in the model's input to override its original instructions or ethical constraints.





B14. Teach your friends about cybersecurity topic of your choice.

During the in-class lab activity 2 about privacy (week 3), I had to teach my group about Multiple Party Computation (MPC) since some of them were still confused about the method. I surprisingly had to teach this twice since I attended the Monday and Thursday lab during that

week (went to the Thursday one to find a group member, which I eventually found). Using the slides as my point of reference, I explained how each person would have their own number they want to keep private, similar to the hospital's raw patient counts. Then they would have to keep a split their number into multiple shares, one share is their own portion to keep, while the other share will be sent to the other people in the group/hospital. The remainder can be shared equally (like Hospital A sending both 30 to B and C) or unequally (like Hospital B and C sending 20 and 15, and 50 and 60). Then each person will add their own kept share with the shares they received from the others, and then the sum they calculated will be shared with others to get the total of everyone's number.

Multiple Party Computation (MPC)

- Three hospitals want to compute the **total number of infected patients** across all hospitals.
- They **cannot share raw patient counts** due to privacy laws (e.g., HIPAA, GDPR).
 - Hospital A: **125** infected patients
 - Hospital B: **98** infected patients
 - Hospital C: **143** infected patients
- Total is $125 + 98 + 143 = 366$.
- Secret sharing technique**

Multiple Party Computation (MPC)

- Each hospital splits its number into **3 shares** (keep one, send two).

Hospital	Share Kept	Sent to A	Sent to B	Sent to C
A (125)	65	—	30	30
B (98)	63	20	—	15
C (143)	33	50	60	—

Multiple Party Computation (MPC)

- Each hospital adds:
 - its **own kept share**
 - Shares received from others

Hospital	Own Share	From Others	Local Sum
A	65	$20 (B) + 50 (C)$	135
B	63	$30 (A) + 60 (C)$	153
C	33	$30 (A) + 15 (B)$	78

Multiple Party Computation (MPC)

- Reveal only **local computed sum**.
 - $135 + 153 + 78 = 366$

Hospital	Revealed Local Sum
A	135
B	153
C	78

B15. Teach an elderly person about cybersecurity topic of your choice.

My grandmother is one of the tech-illiterate elderly who struggles to use and understand technology. She can use it well enough to contact other people and my family, but she is still very much susceptible to the scams and phishing online, especially from random numbers suddenly texting her pretending to be either me or my sister. The image below shows a screenshot of this happening to my grandmother recently, where a random person pretending to be my sister contacted her through WhatsApp (her main media of communication) telling her about buying gold while the current price at that time was still low.

I explained to her that sometimes there can be people/scammers who deceive/trick others into revealing sensitive information (in this case, attempting to steal bank details or get a transfer of money), usually by pretending to be trustworthy entities, such as the situation mentioned, where the attacker pretended to be my sister, someone who obviously my grandmother would trust. I mentioned how this was called “spear phishing”, since the message was targeted towards her and is personalized, as it uses a known relationship. Using the message, I pointed out how she could spot signs of phishing messages:

- The number used was random, and not saved in her contacts (which she usually does for family members). Double check the number, if possible, try to find our actual numbers first and then contact to confirm the information. If not, ask them a personal question that only we would know.
- The profile picture of the number is definitely not my sister, so check the profile picture
- It was a request for something unusual that neither of us (me or my sister) would do, so it's probably not us
- The language they used is slightly off and not how we would usually communicate (we would use more abbreviations and more stickers than emojis)
- They said that this was a “new number”, which is something scammers commonly used to bypass not being in the saved contacts

2:28 PM

50

← +62 817-6405-153



Ini no dari siapa ? 11:32 AM ✓✓

Whose number is this?

No cece 11:33 AM

(This is) my (sister's)
number

Jbeli antam ? 11:33 AM ✓✓

No lama masih aktif 11:33 AM

Did you ended up
buying the gold?

Tapi jarang dibawa 11:33 AM

You
Jbeli antam ?

My old number is still
active, but I rarely
bring it

Iya 😊 11:33 AM

Mumpung harga saat ini lagi turun

Yes. Might as well buy
it while the price is
down

Harga berapaan ? 11:33 AM ✓✓

What's the price?



Delete chat

Unblock

B16. Survey the current state-of-the-art solutions in cybersecurity.

Sources:

- <https://www.sciencedirect.com/science/article/pii/S1877050924013668>
- <https://www.cisco.com/site/us/en/learn/topics/security/what-is-a-next-generation-firewall.html>
- <https://www.cloudflare.com/learning/security/what-is-next-generation-firewall-ngfw/>
- <https://www.enisa.europa.eu/news/enisa-news/what-is-state-of-the-art-in-it-security>
- https://www.teletrust.de/fileadmin/user_upload/2025-09_TeleTrusT_Guideline_State_of_the_art_in_IT_security_EN.pdf

According to guidance from organizations such as ENISA and TeleTrusT, “State of the Art” (SOTA) in IT security is not a fixed technology or single product category, but rather the “best-performing IT security measure available on the market that achieves the legal IT security goal,” where “top performance” means efficiency and effectiveness of safety measures. Depending on the need for protection and the threat, there can be many solutions that can be counted as SOTA. It reflects what is technically “mature” enough for real-world use while still addressing modern attack patterns, regulatory expectations, and operational constraints. In this sense, SOTA is inherently dynamic and evolves alongside changes in attacker techniques, system architectures, and organizational needs.

Within this framing, Next-Generation Firewalls (NGFWs) represent an evolution of traditional firewalls designed specifically to address modern network threats that increasingly operate at the application layer. As described in Cisco’s definition, NGFWs extend beyond basic packet filtering and stateful inspection by incorporating deep packet inspection, application awareness and control, and integrated intrusion prevention capabilities. Rather than simply controlling traffic based on IP addresses or ports, they are able to identify and control specific applications and detect malicious behaviour such as advanced malware and application-layer attacks within allowed traffic streams. This shift is significant because contemporary attacks frequently exploit legitimate protocols such as HTTPS or embed malicious payloads within normal application traffic, making traditional firewall approaches insufficient.

NGFWs are considered SOTA primarily because they align closely with the current evolution of cyber threats and defensive requirements. Modern cybersecurity guidance emphasizes the need for security controls that operate beyond static rule sets and instead provide contextual, behaviour-aware protection. NGFWs achieve this by inspecting traffic at higher OSI layers, particularly the application layer, enabling them to detect threats that would otherwise bypass lower-layer controls. In addition, they integrate real-time threat intelligence feeds, allowing them to dynamically update blocking rules based on emerging global attack patterns. This combination of deep inspection, application-level visibility, and continuous intelligence updates reflects a key characteristic of SOTA systems: adaptability to rapidly changing threats.

Another reason NGFWs are considered SOTA is their role in consolidating multiple security functions into a single platform. Instead of relying on separate systems for firewalling, intrusion detection, malware filtering, and application control, NGFWs unify these capabilities into one integrated enforcement point. This consolidation reduces blind spots between tools, simplifies policy management, and improves response times. From a state-of-the-art perspective, this reflects a broader trend in cybersecurity toward converged, platform-based security architectures that improve operational efficiency while maintaining strong protection.

However, while NGFWs are widely recognized as a mature and advanced security control, their classification as SOTA is contextual rather than absolute. As highlighted in ENISA-style evaluations of security maturity, state-of-the-art does not imply perfection or finality. NGFWs still face limitations, particularly in environments with heavy encryption, where deep inspection may require additional decryption mechanisms, as well as challenges related to performance overhead and complexity in large-scale deployments. For this reason, NGFWs are best understood as a current foundational SOTA technology for network security rather than a complete solution on their own.

Overall, NGFWs represent a state-of-the-art solution because they operationalize modern security principles—application awareness, behavioural analysis, deep packet inspection, and threat intelligence integration—into a deployable and widely adopted system. They bridge the gap between traditional perimeter defense and modern adaptive security needs, making them a central component of contemporary cybersecurity architectures, even as newer paradigms such as Zero Trust and SASE continue to extend beyond them.

B18. Contribute to an open-source project related to cybersecurity.

The open-source project I picked was [CyberChef](#), which is a web application for carrying out cybersecurity operations such as decryptions, encryptions, and many more. Since I'm still starting out in Github and contributions and such, I decided to just do a small documentation improvement in my local fork of the project. I added some missing input/output examples to the operation ToBase32 and FromBase32 since they weren't there originally. The commit to the changes I made can be found [here](#).

All evidence of this, including the changes made in my Docker build of the app, can be found under “\Evidence\B18”.

B19. Find and fix a vulnerability from a GitHub project.

Last year I joined a CFC event/workshop and made a Github project (<https://github.com/Kelly-VI/Kelly-VI.github.io>) deploying my (kinda unfinished) portfolio website as I learned how to use Github, HTML, CSS, and JavaScript (<https://kelly-vl.github.io/>). Scanning my website (<https://app.pentest-tools.com/>) shows that it has a vulnerability of Missing Content-Security-

Policy security header, which is a security measure that prevents Cross-Site Scripting (XSS) vulnerabilities as the header instructs web browsers to enforce specific security rules. If this header is not added, then my website would be easily exploitable by attackers via XSS. Referencing [Mozilla's Content-Security-Policy page](#), I decided to add the following meta tag:


```
<meta
    http-equiv="Content-Security-Policy"
    content="default-src 'self'; script-src 'self'; style-src 'self'; img-
src 'self' data:; media-src 'self';"
/>
```

Which adds all the fetch directives (default-src, script-src, style-src, img-src, media-src), which controls the location from which certain resource types may be loaded, and its specified source expression value, indicating valid sources for that resource type, which in this case is `self`, so resources of the given type may only be loaded from the same origin as the document/website. Then I added the change, committed, and pushed it to Github in this [commit](#).

B20. Enhance the security of a GitHub project.

My other unit, CITS3403: Agile Web Development, has a group project where we have to create a web application. One of the requirements of this project is for users to be able to login and logout, which we have created in login and register forms. When logged in, this creates a session cookie. Originally, our configuration settings did not have anything that dealt with cookie security, so I added some using the [Flask Configuration Hardening](#) website as my guide. The added configuration values are as follows:

- SESSION_COOKIE_HTTPONLY: Browser will not allow JavaScript access to cookies marked as "HTTP only"
- SESSION_COOKIE_SECURE: Browser will only send cookies over HTTPS in production (True). Since our project is currently in development/debugging, it will be false (based on if condition)
- SESSION_COOKIE_SAMESITE: Restrict how cookies are sent with requests from external sites. ['Lax'](#) was the recommended option as it prevents sending cookies with CSRF-prone requests from external sites



A screenshot of a code editor window titled 'Agile-Web-Group-Project'. The editor shows a Python file named 'config.py' with the following code:

```
1 from dotenv import load_dotenv
2 import os
3 from datetime import timedelta
4
5 load_dotenv()
6
7 class Config:
8     # Secret key for session logic
9     SECRET_KEY = os.getenv('SECRET_KEY')
10
11     # Flag error if missing secret key environment var:
12     if not SECRET_KEY:
13         raise Exception("SECRET_KEY is not defined, check if you have the .env file!")
14     # Configure the database URI and initialize the database
15     SQLALCHEMY_DATABASE_URI = 'sqlite:///app.db'
16
17     # Configure security settings for cookies/sessions
18     SESSION_COOKIE_HTTPONLY = True # browser will not allow JavaScript access to cookies marked as "HTTP only"
19     SESSION_COOKIE_SECURE = os.getenv('FLASK_ENV') == 'production' # only send cookies over HTTPS in production (True), for now in debug:
20     SESSION_COOKIE_SAMESITE = 'Lax' # prevents sending cookies with CSRF-prone requests from external sites (e.g. submitting a form)
21
22
23     # It's generally bad practice to hardcode a key even on the chance that environment variable is inaccessible.
24     # If a SECRET_KEY can't be retrieved from env, the app should raise an error instead.
```

B21. Design and implement a cybersecurity learning activity.

I decided to make a simple Kahoot quiz on cybersecurity awareness. It's a mix of introduction cybersecurity topics such as phishing, authentication, encryption, network security, and incident response thinking. There's 15 questions in total.

<https://create.kahoot.it/share/cybersecurity-learning-activity/2bbe946a-80fe-4001-83cf-de191abd59bc>



B22. Enhance the cybersecurity of a website from your community.


Security Headers
 by snyk

Scan your site now

☐ Hide results
 ☒ Follow redirects

Security Report Summary




Site:	https://acetutoring.com.au/
IP Address:	85.187.128.7
Report Time:	02 May 2026 15:22:11 UTC
Headers:	✔ Strict-Transport-Security ✔ X-Frame-Options ✔ X-Content-Type-Options ✖ Content-Security-Policy ✖ Referrer-Policy ✖ Permissions-Policy
Advanced:	Not bad... Maybe you should perform a deeper security analysis of your website and APIs: Try Now

Missing Headers

Content-Security-Policy	Content Security Policy is an effective measure to protect your site from XSS attacks. By whitelisting sources of approved content, you can prevent the browser from loading malicious assets.
Referrer-Policy	Referrer Policy is a new header that allows a site to control how much information the browser includes with navigations away from a document and should be set by all sites.
Permissions-Policy	Permissions Policy is a new header that allows a site to control which features and APIs can be used in the browser.

My community would be my tutoring centre. A scan (<https://securityheaders.io/>) of their website shows that one of the vulnerabilities they have is a Missing Referrer-Policy header, which is the header that controls how much referrer information the browser will send with each request originated from the current website. Essentially, this means that if a user visits the tutoring centre’s website and clicks on a link from that page to another page, then the browser will send the originating URL in the Referrer header to the destination website. This can be used for user tracking and may lead to sensitive information being leaked for the user. Although I do not have the authority to actually fix the website, I would fix this issue by adding a Referrer-Policy directive to every HTTP response in the web server (<https://scotthelme.co.uk/a-new-security-header-referrer-policy/>).

The most widely used and default directive is `strict-origin-when-cross-origin` (https://developer.mozilla.org/en-US/docs/Web/Security/Practical_implementation_guides/Referrer_policy) as it sends the full Referrer header information on same-origin requests, only the URL without the path (origin) on cross-origin HTTP requests, and nothing for downgrading to HTTP. Inspecting the website’s HTTP Response Headers using the browser’s DevTools (under Network tab) reveals that the server they use is [LiteSpeed](#). According to their [documentation](#), the security header can be added via WebAdmin Console (navigate to Web Admin > Configurations > Your Virtual Hosts > Context, click the Add button to add the headers, then save), or added to .htaccess file in the web application’s root directory, adding the line `Header set Referrer-Policy "strict-origin-when-cross-origin"`.



B23. Test an intrusion detection system and discuss its effectiveness.

Sources:

- <https://www.youtube.com/watch?v=W1pb9DFCXLw&list=PLpPXZRVU-dX33VNUeqWrMmBNf5FeKVmi->
- <https://www.youtube.com/watch?v=PYP0YH2PVuo>
- <https://en.wikipedia.org/wiki/EternalBlue>
- <https://learn.microsoft.com/en-us/security-updates/Securitybulletins/2017/ms17-010>
- <https://www.cve.org/CVERecord?id=CVE-2017-0144>
- <https://learn.microsoft.com/en-us/troubleshoot/windows-server/networking/inter-process-communication-share-null-session>
- <https://talk-about-it.ca/cant-connect-admin-ipc-shares-computer/>

Snort is an open-source IPS that can do real-time traffic analysis and packet logging. They can detect and then drop/neutralize malicious traffic before it comes into the network. Packets are processed by Snort by first sending them through the Snort Decoder, which determine which underlying protocol is used in the packet (e.g. Ethernet, IP, TCP) by looking through the fields in the packet's error, while also looking out for any errors or anomalies. Then they are sent to Preprocessors, which are plugins designed for Snort to handle certain types of traffic by standardizing/normalizing the format of the data/packet into something Snort can use. After that, the packet is sent to through Snort's route, where rules defined by either Talos (the company maintaining Snort) or custom rules by the user can flag traffic that is either detected as malicious or against policy. If it is, then Snort will drop the packet completely from the

network so that it doesn't arrive to the destination. All of these processes will be logged by Snort for users to view and analyze.

For my test, I decided to follow through the Cisco Talos' Snort3 Tutorial via the installation of [Snort3 Docker Container](#). I first run the command `snort -c /home/snorty/snort3/etc/snort/snort.lua` to specify which configuration file I will be using (`snort.lua`)

Terminal

```
snorty@snort3:~$ snort -c /home/snorty/snort3/etc/snort/snort.lua
-----
o")~  Snort++ 3.10.0.0
-----
Loading /home/snorty/snort3/etc/snort/snort.lua:
Loading snort_defaults.lua:
Finished snort_defaults.lua:
    output
    so_proxy
    classifications
    references
    binder
    wizard
    appid
    js_norm
```

Terminal

```
Loading file_id.rules_file:
Loading file_magic.rules:
Finished file_magic.rules:
Finished file_id.rules_file:
-----
ips policies rule stats
      id loaded  shared enabled   file
      0   219     0     219  /home/snorty/snort3/etc/snort/snort.lua
-----
rule counts
  total rules loaded: 219
    text rules: 219
  option chains: 219
  chain headers: 1
-----
      memory scale: KB
      total memory: 71.2812
    pattern memory: 19.6484
  match list memory: 28.4375
    transition memory: 22.9453
appid: MaxRss diff: 3200
appid: patterns loaded: 300
-----
pcap DAQ configured to passive.

Snort successfully validated the configuration (with 0 warnings).
```

`snort -R /home/snorty/snort3/etc/snort/samples.rules -r <pcap name> -A alert\_test -n 100000` is the command to run Snort in IDS mode against a pcap.

Lab 1 of the Tutorial (basically learning how to run snort) makes us run `./run.sh`, which runs the command `snort -q --talos -r get.pcap`, which runs Snort in quiet mode (not showing Snort banner or status report, `-q`), uses talos tweak and alert talos logger (easy way to run an inline rules test against a local rules file as well as getting an easy to read output, `--talos`), and running this against the `get.pcap` file. The alert talos logger splits alerts up based on our (local rule) pcap that caused the alert.

```
snorty@snort3:~/examples/intro/lab1$ ./run.sh
running:
snort -q --talos -r get.pcap

##### get.pcap #####
      [1:1:0] Hello World (alerts: 1)
#####
-----
rule profile (all, sorted by total_time)
#      gid    sid rev    checks matches alerts time (us) avg/check avg/match avg/non-match timeouts suspends rule_time (%)
=      ==     ==  ==    =====
=====
```

The rule that triggered the alert was this:

```
snorty@snort3:~/examples/intro/lab1$ cat local.rules
alert http (
    msg:"Hello World";
    flow:to_server,established;
    http_uri;
    content:"/hello",nocase;
    sid:1;
)
```

Lab 2 (which will be my main test of Snort) makes us run `snort -q --talos --rule-path ~/snort3/etc/rules/ -r eternalblue.pcap`. This command uses the free community rules file that is already installed in the Docker, and running it against `eternalblue.pcap` file to let Snort analyze and detect traffic matching known attack signatures related to the EternalBlue exploit, which was an exploit based on a zero-day vulnerability in Microsoft Windows software that allowed users to gain access to any number of computers connected to a network and exploits Microsoft's implementation of the Server Message Block (SMB) protocol. The SMBv1 server in various versions of Microsoft Windows mishandles specially crafted packets from remote attackers, allowing them to remotely execute arbitrary code on the target computer.

```
snorty@snort3:~/examples/intro/lab2$ snort -q --talos --rule-path ~/snort3/etc/rules/ -r eternalblue.pcap

##### eternalblue.pcap #####
[1:41978:5] OS-WINDOWS Microsoft Windows SMB remote code execution attempt (alerts: 27)
[1:42331:5] MALWARE-CNC Win.Trojan.Doublepulsar variant process injection command (alerts: 3)
[1:42340:4] OS-WINDOWS Microsoft Windows SMB anonymous session IPC share access attempt (alerts: 3)
[1:42944:2] OS-WINDOWS Microsoft Windows SMB remote code execution attempt (alerts: 3)
#####

-----
rule profile (all, sorted by total_time)
#      gid      sid rev      checks matches alerts time (us) avg/check avg/match avg/non-match timeouts suspends rule_time (%)
=      ==      == ==      =====
1      1      3041  5      396      0      0      1515      3      0      3      0      0      1.72638
2      1      3033  9      396      0      0      1515      3      0      3      0      0      1.72638
3      1      3025  9      396      0      0      1515      3      0      3      0      0      1.72638
4      1      3031  9      396      0      0      1507      3      0      3      0      0      1.71726
5      1      3023  9      396      0      0      1507      3      0      3      0      0      1.71726
6      1      3039  9      396      0      0      1507      3      0      3      0      0      1.71726
7      1      3138  11     396      0      0      1505      3      0      3      0      0      1.71498
```

Alert 1 ([1:41978:5]) detected `OS-WINDOWS Microsoft Windows SMB remote code execution attempt`, which is likely a detection of the ExternalBlue SMB exploit. This triggered 27 alerts, which meant that 27 packets matched the rules. Alert 4 [1:42944:2] also detected the same SMB-based remote code execution activity, but this time with a different signature ID than the first alert (42944 vs. 41978) along with triggering 3 alerts, indicating 3 packets matching the signature. The different Snort rules detect different stages, variants, or traffic patterns of SMB exploitation, hence why multiple rules are triggered from this same pcap. Additionally, Alert 3 [1:42944:2] detected `OS-WINDOWS Microsoft Windows SMB anonymous session IPC (Inter-Process Communication) share access attempt`, which indicates that a system attempted to access a Windows SMB IPC share anonymously without normal authentication (*IPC share is a special hidden administrative share used internally by Windows networking services to manage the computer environment or network). Snort might be detecting this because anonymous SMB access is now uncommon in modern Windows Server versions, so it might be treating this signature as suspicious behaviours in case attackers are trying to do reconnaissance in preparation for their EternalBlue exploitation.



Based on the output, Snort has successfully detected multiple indicators of malicious SMB activity within the `eternalblue.pcap` capture as the IDS generated alerts for MS Windows SMB remote code execution attempts, which indicates that Snort's Talos ruleset was able to identify several stages of the EternalBlue attack chain and correctly recognize the traffic patterns associated with it. The high number of alerts, especially from Alert 1, suggests that the IDS consistently identified repeated exploit-related packets throughout the capture. Overall, this means that Snort is quite effective at detecting known attacks using signature-based rules, identifying multiple stages of an attack cycle, and analyzing recorded network traffic. Since the same detection rule engine applies to both recorded and live network traffic, Snort would also be quite effective in processing packets in real time and then dropping the malicious packets that triggered the alerts, keeping the system safe.

However, since Snort is rule-based, it is limited by its predefined rules and signatures of attacks. So if, for example, the EternalBlue exploit had been initially deployed into a system when the vulnerability wasn't discovered yet (zero-day), then the attack might have gone undetected since there would've been no rules for the "new" attack yet.


Overall, Snort is quite effective at detecting attacks and suspicious activity when comparing it with its defined ruleset.

B24. Design and implement access control of your choice.


I have a discord server with my friends where we do random things (as with all friend groups). I decided to mess around in the permissions settings of the server since Discord is one of the easiest ways to implement Role-Based Access Control (RBAC) directly through server roles and permissions. One of the roles is administrator, which is the role that allows full server control, so for example, administrator would have the permissions to Manage Channels (allows members to create, edit, or delete channels), Manage Server (allow members to change this server's name, switch regions, view all invites...), Kick, Approve, and Reject Members, Ban Members, etc. Then I created another role "Thou art guest" where they cannot do any of the permissions mentioned above and can only access general channels. I also created some private channels where only administrators can view it and others cannot, since it was specifically made for administrators only. So if I try to view the server as the "Thou art guest" role, I would not see these private channels.



Administrator POV:




Guest POV:



B25. Design and implement a threat intelligence module of your choice.

B26. Help another student in this unit struggling to understand/learn a cybersecurity concept.



The screenshot shows a web page for CITS2006. The header includes a menu icon, a logo, and the text 'CITS2006'. On the right, there is a search bar with the text 'Search...' and a 'Ctrl K' shortcut. The main content area has a title '1.2.4. One-way vs. collision-free properties'. Below the title, there is a paragraph explaining that one-way hash functions are not necessarily collision-free, but collision-free means it is impossible to find two different messages with the same hash value. It also states that one-way hash functions are designed to be computationally infeasible to find a message with a specific hash value. This means it is possible to find two different messages with the same hash value, but it is computationally infeasible to find such messages. Below this paragraph, it says 'You are given a oneway.py file:' and provides a code block with the command: `wget https://github.com/uwacyber/cits2006/raw/2026/cits2006-labs/files/oneway.py`. At the bottom, it says 'Open and inspect the code. It tests a one-way property of hash functions. It currently implements MD5, but you can easily replace that with other hash functions to test.'

My friend and I met up on Wednesday, 4 March 2026 to do Lab 1 together. While going through the task, we stumbled upon this section of Lab 1 which gave us a bit of confusion. After researching a bit more about the property and discussing it together, we came up with this understanding:

- One-way property means that it is easy to get the hash of the message (forward direction), but it is very difficult/computationally infeasible to do the inverse, i.e. figure out the message from its hash. Thus, even if two messages have the same hash value, it is very difficult to find what the messages are that gave those same hash results.
- Collision-free property (which is apparently also known as collision resistance property) means that it is hard to find two inputs that hash to the same output value. However, just because it is hard (computationally infeasible), doesn't mean that it doesn't exist or it is impossible to find one since the Pigeonhole Principle states that any hash function that has more inputs than outputs will have collisions. This is because the input can technically have infinite bits, while the output is limited to a fixed bit size, such as 128-bit hash value for MD5 and 256-bit for SHA-256.
- So even in one-way property, collisions do exist, but they are very hard to find/compute intentionally.

B27. Apply a learned concept in this unit to a real-world application/problem/environment.


- <https://keepass.info/help/base/security.html>

I've always wanted to know how cryptography actually works in storing passwords, so I decided to do a deep dive into the password manager that I am using, KeePass, which is a real-world application. In the lectures, we learned about cryptography, encryption (protect data by converting it into unreadable ciphertext which is decrypted with a key), and hashing (one-way transformation of data used to verify integrity and secure password authentication), and that storing the plaintext of passwords is bad and that we should hash and encrypt them so that even if the password is stolen, the attackers can't read them. KeePass manages and stores the user's passwords in a database. These database files are encrypted (the whole database is encrypted, not just the password. This includes usernames, URLs, notes, etc). Since I have the 2.x version of KeePass, it uses the AES-256 and ChaCha20, which are well-known and thoroughly analyzed algorithms that are considered to be very secure. Every time a database is saved, an initialization vector (IV) is generated randomly, so even if there were multiple databases encrypted with the same master key (for example, for backups), the database will still be secure. This encryption ensures confidentiality of all the stored credentials in the database and thus protects sensitive data even in case of system compromise.

Additionally, SHA-256 is also used for compressing the components of the master key (consisting of a master password, a key file and/or a key that is protected using the current Windows user account). To open a database file, all components of the master key are required, so if one of the components are lost/forgotten, then all data stored in the database is lost as there is no backdoor nor universal key that can open the database. This is because SHA-256 is a cryptographic hash function that generates a 256-bit signature from input data, and hashing (as we have learned) is one-way and not reversible, so any missing input data cannot be reconstructed from the hash and changing one part of the input will produce a completely different hash.

B28. Produce a cyber safety flyer for (choose 1): elders, high school students, CEOs, Uni students.

I made a cyber safety flyer for university students, specifically how they would use similar passwords every time the university asks them to change their passwords every semester (I too am unfortunately guilty for this). Changing and coming up with new passwords every time is a hassle for students, especially during the UWA Data Breach 2025, where students had to keep changing their passwords multiple times in the span of a week, and I'm assuming most used the same password they have been using previously, with minor variations such as a number change, or adding one extra character, etc. Hence why I decided to make this flyer using Canva (used a template for the background).



HEY YOU! do YOU:

Reuse the **SAME** password (with minor variation) **EVERY TIME** they ask you to change the password every semester?

don't!

BAD STUDENT HABIT!

That makes it easier for your account to get hacked!!!

USE: PASSWORD MANAGERS

Stored and manages all your **STRONG** and **UNIQUE** passwords for you

Examples:

- Bitwarden
- Proton Pass

START USING THEM NOW!

B29. Find a CVE in this year and fix it using three different generative AI systems (e.g., ChatGPT, Gemini), comparing the consistency.

The CVE I found was [CVE-2026-20128](#), which is about Cisco Catalyst SD-WAN Manager having a vulnerability in the Data Collection Agent (DCA) that allowed unauthenticated, remote attackers to gain DCA user privileges on an affected system due to the presence of a credential file for the DCA user, which can be exploited via a crafted HTTP request to read the file containing the DCA password from the affected system. This was a high severity vulnerability (according to the Common Vulnerability Scoring System, CVSS) and is related to [CWE-257: Storing Passwords in a Recoverable Format](#).

I will fix this CVE using ChatGPT, Gemini, and Copilot. All evidence of the prompt and then given output from each genAI system will be under the directory “B29” in Evidence.

The comparisons of the consistency are as follows:

- All models mention to upgrade to the software version 20.18+ that has patched the vulnerability. Copilot specifically gives the table of fixed software releases and mention the versions that contain the fixes to the vulnerability.
- All models also give an “alternative” to not being able to upgrade immediately, i.e. applying “hardening measures”, although ChatGPT doesn’t explicitly say “hardening measures”, opting to simply list them out while including “If immediate upgrade isn’t possible” or “Reduce exposure until patched”
- ChatGPT and Copilot each gives a “bottom line” or “summary” at the end of their output, while Gemini doesn’t
- Copilot likes to reference from the [actual source](#) of the announcement (including the tables for all the branch-specific versions and code and all) and other official sources (such as other [Cisco websites](#) for Remediation or [Cyber.gov.au](#)) for additional defensive steps, including how it emphasizes that patching is the only fix and how there are no workarounds for the vulnerability. Most of their response is based on the official sources/guides, making their response very specific and detailed, almost like a summary of the sources. ChatGPT and Gemini simply give brief/generic/simplified mitigations that may or may not work.
 - o Since Cisco has officially stated that there is no workaround for the vulnerability, the only “real” fix is to essentially wait for the patch, which Copilot emphasizes. ChatGPT and Gemini attempted to give other “viable” mitigations, but these mitigations are not that effective or cannot be meaningfully mitigated without patching, so they are kind of contradicting themselves with their mitigations.
 - o Gemini has more assumptions/unverified claims in their response. While Cisco has released a patch to fix the vulnerability, they did not state what they did to fix it. However, Gemini states that “Versions 20.18+ have removed the insecure credential handling for DCA” and “patched the directory traversal/file read

mechanism used to access it” when Cisco has made no mentions of this whatsoever.

- Overall, Copilot is the most consistent as it references official sources about the problem and how to mitigate them. Gemini is the least consistent, having a lot more assumptions and unverified claims in its response. ChatGPT sits somewhere in between both, but is definitely less consistent than Copilot.